

SVD Factorization for Tall-and-Fat Matrices on Map/Reduce Architectures

Burak Bayramlı

October 15, 2013

Abstract

We demonstrate an implementation for an approximate rank- k SVD factorization, combining well-known randomized projection techniques with previously implemented map/reduce solutions in order to compute steps of the random projection based SVD procedure, such as QR and SVD. We structure the problem in a way that it reduces to Cholesky and SVD factorizations on $k \times k$ matrices computed on a single machine, greatly easing the computability of the problem.

1 Introduction

[1] presents many excellent techniques for utilizing map/reduce architectures to compute QR and SVD for the so-called tall-and-skinny matrices. QR factorization is turned into an $A^T A$ computation problem to be computed in parallel using map/reduce, and its key element the Cholesky decomposition, can be performed on a single machine. Let's use $C = A^T A$ and, since

$$C = A^T A = (QR)^T (QR) = R^T Q^T QR = R^T R$$

and because Cholesky factorization of an $n \times n$ symmetric positive definite matrix is

$$C = LL^T$$

where L is an $n \times n$ lower triangular matrix, and R is upper triangular, we can conclude if we factorize C into L and L^T , this implies $C = LL^T = RR^T$, we have a method of calculating R of QR using Cholesky factorization

on $A^T A$. The key observation here is $A^T A$ computation results an $n \times n$ matrix and if A is “skinny” then n is relatively small (in the thousands), then Cholesky decomposition can be executed on a small $n \times n$ matrix on a single computer utilizing an already available function in a scientific computing library. Q is computed simply as $Q = AR^{-1}$. This again is relatively cheap because R is $n \times n$, the inverse is computed locally, matrix multiplication with A can be performed through map/reduce.

SVD is an additional step. SVD decomposition is

$$A = U\Sigma V^T$$

If we expand it with $A = QR$

$$QR = U\Sigma V^T$$

$$R = Q^T U \Sigma V^T$$

Let’s call $\tilde{U} = Q^T U$

$$R = \tilde{U} \Sigma V^T$$

This means if we run a local SVD on R (we just calculated above with Cholesky) which is an $n \times n$ matrix, we will have calculated \tilde{U} , the real Σ , and real V^T .

Now we have a map/reduce way of calculating QR and SVD on $m \times n$ matrices where n is small.

1.1 Approximate rank-k SVD

Switching gears, we look at another method for calculating SVD. The motivation is while computing SVD, if n is large, creating a “fat” matrix which might have columns in the billions would require reducing the dimensionality of the problem. According to [2], one way to achieve is through random projection. First we draw an $n \times k$ Gaussian random matrix Ω . Then we calculate

$$Y = A\Omega$$

We perform QR decomposition on Y

$$Y = QR$$

Then form $k \times n$ matrix

$$B = Q^T A$$

Then we can calculate SVD on this small matrix

$$B = \hat{U} \Sigma V^T$$

Then form the matrix

$$U = Q \hat{U}$$

The main idea is based on

$$A = Q Q^T A$$

if replace Q which comes from random projection Y ,

$$A \approx \tilde{Q} \tilde{Q}^T A$$

Q and R of the projection are close to that of A . In the multiplication above R is called B where $B = \tilde{Q}^T A$, and,

$$A \approx \tilde{Q} B$$

then, as in [1], we can take SVD of B and apply the same transition rules to obtain an approximate U of A .

This approximation works because of the fact that projecting points to a random subspace preserves distances between points, or in detail, projecting the n -point subset onto a random subspace of $O(\log n / \epsilon^2)$ dimensions only changes the interpoint distances by $(1 \pm \epsilon)$ with positive probability [3]. It is also said that Y is a good representation of the span of A .

1.2 Combining Both Methods

Our idea was using approximate k -rank SVD calculation steps where $k \ll n$, and using map/reduce based QR and SVD methods to implement those steps. By utilizing random projection, we would be able to work in a smaller dimension which would translate to local Cholesky, and SVD calls on $k \times k$ matrices that can be performed in a speedy manner. Below we outline each map/reduce job.

First random projection job (whose reduce is a no-op). Each value of A will arrive to the algorithm as a key and value pair. Key is line number or

```

random_projection_map(key, value)
1   input A
2   returns Y
3   Tokenize value and pick out id value pairs
4   result = zeros(1,k)
5   for each  $j^{\text{th}}$  token  $\in$  value
6       Initialize seed with j
7       j = generate k random numbers
8       result = result + r · token[j]
9   emit key, result

```

```

ATAcholesky_job_map(key k, value a)
1   for i, row in enumerate aT
2       emit i, row

```

other identifier per row of A . Value is a collection of id value pairs where id is column id this time, and value is the value for that column. Sparsity is handled through this format, if an id for a column does not appear in a row of A , it is assumed to be zero. The resulting Y matrix has dimensions $m \times k$.

The `cholesky_job_final_local_reduce` step is a function provided in most map/reduce frameworks, it is a central point that collects the output of all reducers, naturally a single machine which makes it ideal to execute the final Cholesky call on by now a very small ($k \times k$) matrix. The output is R .

There is no reducer in the `Q_job`, it is a very simple procedure, it merely computes multiplication between row of Y and a local matrix R . Matrix R is very small, $k \times k$, hence it can be kept locally in every node. The initialization is used to store the inverse of R locally, once the mapper is initialized, it will always use the same R^{-1} for every multiplication.

The job above takes an AQ matrix which is assumed to be a join between A and Q , per row, based on key. We split the row and deduce the A part and the Q part, then iterate cells of A one by one, which is assumed to

```

cholesky_job_reduce(key, value)
1   emit k, sum( $v_j^k$ )

```

```

    cholesky_job_final_local_reduce(key, value)
1   result = cholesky( $A_{\text{sum}}$ )
2   emit result

```

```

    Q_job_map(key, value)
1   During initialization,  $R_{\text{inv}} = R^{-1}$ , store it once for each mapper
2   for row in Y
3       emit key, row ·  $R_{\text{inv}}$ 

```

```

     $A^T Q$ _job_map(key, value)
1   left = row from A
2   right = row from Q
3   for each non-zero  $j^{\text{th}}$  cell in left
4       emit j, left[j] · right

```

```

     $A^T Q$ _job_reduce(key, value)
1   returns  $B^T$ 
2   result = zeros(1, k)
3   for row in value
4       result = result + row
5   emit key, result

```

```

Q $\tilde{U}$ _job_map(key,value)
1   input Q, R
2   returns U
3   initialization  $\tilde{U}$  = svd of R
4   for row in Q
5       emit key, row ·  $\tilde{U}$ 

```

```

map_reduce_svd
1   Y = random_projection_map(A)
2    $R_Y = A^T A_{\text{cholesky\_job}}(Y)$ 
3    $Q_Y = Q_{\text{job}}$ 
4    $R_{BT} = A^T A_{\text{cholesky\_job}}(B^T)$ 
5    $U = Q_{\tilde{U}\text{-job}}(R_{BT}, Q)$ 

```

be sparse, and multiply the entire row of Q . Then for each j^{th} non-zero cell of A , we multiply this value with the row from Q and emit the multiplication result with key j .

The $Q^T A$ job's formula can be seen at 1.1. For implementation purposes we changed this formula into

$$B^T = A^T Q$$

because as output we needed to have a $n \times k$ matrix instead of a $k \times n$ one, which would allow us to use map/reduce SVD that translates into a local Cholesky and SVD on $k \times k$ matrices. Since we take SVD of B^T instead of B , that changes the output as well,

$$B = U \Sigma V^T$$

becomes

$$B^T = V \Sigma U^T$$

In other words, in order to obtain U of B , we need to take $(U_{BT}^T)^T$ from the SVD of B^T . This is how $A^T A$ Cholesky Job is called, this time with B^T as its input data.

1.3 Discussion

We performed our experiments on the Netflix dataset which has about 100 million from over 480,000 customers on 17770 movies. The implementation was programmed on Sasha distributed framework [5], and SVD calculation on the full dataset with $k = 7$ on two notebook computers, utilizing in total 6 cores took 20 minutes. Scipy SVD calculation on the same dataset is much faster, however, we need to stress our algorithms are prepared for cases where N is very large, i.e. in the billions. As such, for example during projection we did not simply create and pre-store a $n \times k$ random matrix and multiply multiple rows of A with this matrix. This would certainly be possible for Netflix data where n is relatively small, but would not work well in cases where A is “fat”. All code relevant for this paper can be found here [6].

There are only two passes necessary on the full dataset, and three passes on m rows but with reduced k dimensions this time. Perhaps predictably, the procedure spends most of its time at $A^T Q$ Job. This step performs not only a join between A and Q , it also emits k cells per non-zero value of A 's rows, then creates partial sums these k vectors creating $n \times k$ result. If for simplicity we assume k non-zero cells in each A row, the complexity of this step would be $O(mk)$.

References

- [1] Gleich, Benson, Demmel, *Direct QR factorizations for tall-and-skinny matrices in MapReduce architectures*, arXiv:1301.1071 [cs.DC], 2013
- [2] N. Halko, *Randomized methods for computing low-rank approximations of matrices*, University of Colorado, Boulder, 2010
- [3] S. Dangupta, A. Gupta *An Elementary Proof of a Theorem of Johnson and Lindenstrauss*, Wiley Periodicals, 2002
- [4] M. Kurucz, A. A. Benczúr, K. Csalogány, *Methods for large scale SVD with missing values*, ACM, 2007
- [5] B. Bayramli, *Sasha Framework*, `git@github.com:burakbayramli/sasha.git` Github, 2013
- [6] B. Bayramli, *Map/Reduce Code for Netflix SVD Analysis*, `http://github.com/burakbayramli/classnotes/tree/master/stat/stat_mr_rnd_svd/sasha`, Github, 2013